

Andrea Ji Woo NAM SONG

Stanford, CA | +1 (650) 788-9137 | anamsong@stanford.edu | linkedin.com/in/andrea-nam-5862791b8/ | <https://andrea0910.github.io>

EDUCATION

Stanford University

M.S. Candidate, Computer Science with a Concentration in Artificial Intelligence

B.S. Management Science & Engineering with a Concentration in Finance and Decision

• Selected Courses: CS 224R Deep Reinforcement Learning; CS 224N Natural Language Processing with Deep Learning; CS 229 Machine Learning; CS 221 Artificial Intelligence; CS 153 Frontier Systems; CS 246 Mining Massive Data Sets; CS 238 Decision Making Under Uncertainty; MS&E 180: Organization: Theory & Management; ENGR 108: Technology Entrepreneurship

Stanford, CA

Class of 2026

Class of 2025

WORK EXPERIENCES

microl

Human Data Manager (HDM)

• Built and owned performance dashboards monitoring data quality, throughput, and error rates across 100+ AI trainers, giving program leadership real-time visibility into training-data health

• Diagnosed and resolved data-pipeline bottlenecks through targeted QA workflows, improving deliverable accuracy and reducing rework cycles for client reporting

• Authored evaluation rubrics and QA standards adopted program-wide, aligning technical, ops, and client teams on a shared quality bar

Remote

August 2025 - November 2025

Bucketplace (Ohouse)

US Strategy & Operations Intern—Product Launch

• Led U.S. market entry research through 10+ user interviews and evaluation of 100+ SKUs, shaping launch strategy and curating the MVP product catalog directly with Head of America

• Translated and SEO-optimized 100+ Korean product descriptions for U.S. launch, improving discoverability and localized user experience

Remote

May 2025 - July 2025

Banco Santander International - Private Banking

Investment Department Intern

• Built and backtested an r^* neutral-rate model (U.S. and Eurozone, 1980–present); model forecast a U.S. rate cut to 2.5–3% and flagged Eurozone inflation risk via a $2\%+ r^*$ estimate

• Presented model methodology and predictions directly to the directive committee, including the CIO and Director of Investments, to inform scenario analysis and portfolio strategy

Miami, FL

May 2024 - August 2024

Hanwha Life - Global Internship Program

VC Intern - Strategic Investment Operations Team

• Benchmarked 20+ GenAI products (B2B and B2C) across Asia and North America, evaluating UX differentiation to inform AI-product partnership strategy

• Researched U.S. fintech personalization trends and produced investment recommendations calibrated to the Korean market, synthesized into executive reports for the investment committee

Seoul, South Korea

June 2023 - August 2023

PROJECTS

Adaptive Test-Time Sampling for RL-Fine-Tuned Reasoning Policies

Researcher, CS224R: Deep Reinforcement Learning

• Investigating whether reasoning policies trained with on-policy RL (IPO, RLOO) benefit differently from adaptive test-time compute allocation than SFT-only policies on the Countdown math reasoning task (Qwen 2.5 0.5B base)

• Implementing a self-consistency-based stopping rule that halts sampling once the plurality answer reaches an agreement threshold, redirecting compute toward problems where initial generations disagree

• SFT baseline trained and evaluated (88% test token accuracy, ~71% pass@16); IPO and RLOO checkpoints and the adaptive vs. fixed-k comparison at matched compute budgets in progress

Stanford, CA

April 2026 - Present

Evaluation User-Style Adaptation for Professional Text Generation

Researcher, CS224N: Natural Language Processing with Deep Learning

• Extended HyPerAlign, a prompt-based personalization framework, with a structured 11-field email schema augmented with user-distinctive vocabulary automatically extracted from writing samples, plus a lightweight prompt optimizer that identifies the minimal effective style field subset

• A 3-field schema (greeting style, punctuation, distinctive voice) closed >50% of the gap to the performance ceiling with a single exemplar, outperforming both parameter-efficient fine-tuning and unstructured hypothesis prompting

Stanford, CA

January 2026 - March 2026

Stochastic Routing Process for Cross-Border Payments

Co-Researcher, CS 238: Decision Making Under Uncertainty

• Designed a greedy stochastic routing framework over institution-currency networks; Monte Carlo simulations over 10 years of FX data achieved 100% routing success on feasible corridors

• Final amounts systematically exceeded the commercial mid-market benchmark across all tested corridors; model correctly identified structural failure cases where network asymmetries made routing infeasible

Stanford, CA

October 2025 - December 2025

SKILLS & INTERESTS

Programming: Python (PyTorch, HuggingFace Transformers, pandas, NumPy), R, SQL

ML/AI: PyTorch, Hugging Face Transformers, on-policy RL training (IPO, RLOO), supervised fine-tuning, LLM evaluation, Monte Carlo methods

Languages: Spanish (Native), Korean (Native), English (Fluent), Portuguese, Chinese