Andrea Ji Woo NAM SONG
Stanford, CA | +1 (650) 788-9137 | anamsong@stanford.edu | linkedin.com/in/andrea-nam-5862791b8/ | https://andrea0910.github.io

EDUCATION

Stanford University **Stanford, CA**
M.S. Candidate, Computer Science with a Concentration in Artificial Intelligence *Class of 2026*
B.S. Management Science & Engineering with a Concentration in Finance and Decision *Class of 2025*
● Selected Courses: CS 224R Deep Reinforcement Learning; CS 224N Natural Language Processing with Deep Learning; CS 229 Machine Learning; CS 221 Artificial Intelligence; CS 153 Frontier Systems; CS 246 Mining Massive Data Sets; CS 238 Decision Making Under Uncertainty; MS&E 180: Organization: Theory & Management; ENGR 108: Technology Entrepreneurship

WORK EXPERIENCES

micro1 **Remote**
Human Data Manager (HDM) *August 2025 - November 2025*
● Built and owned performance dashboards monitoring data quality, throughput, and error rates across 100+ AI trainers, giving program leadership real-time visibility into training-data health
● Diagnosed and resolved data-pipeline bottlenecks through targeted QA workflows, improving deliverable accuracy and reducing rework cycles for client reporting
● Authored evaluation rubrics and QA standards adopted program-wide, aligning technical, ops, and client teams on a shared quality bar

Bucketplace (Ohouse) **Remote**
US Strategy & Operations Intern—Product Launch *May 2025 - July 2025*
● Led U.S. market entry research through 10+ user interviews and evaluation of 100+ SKUs, shaping launch strategy and curating the MVP product catalog directly with Head of America
● Translated and SEO-optimized 100+ Korean product descriptions for U.S. launch, improving discoverability and localized user experience

Banco Santander International - Private Banking **Miami, FL**
Investment Department Intern *May 2024 - August 2024*
● Built and backtested an r* neutral-rate model (U.S. and Eurozone, 1980–present); model forecast a U.S. rate cut to 2.5–3% and flagged Eurozone inflation risk via a 2%+ r* estimate
● Presented model methodology and predictions directly to the directive committee, including the CIO and Director of Investments, to inform scenario analysis and portfolio strategy

Hanwha Life - Global Internship Program **Seoul, South Korea**
VC Intern - Strategic Investment Operations Team *June 2023 - August 2023*
● Benchmarked 20+ GenAI products (B2B and B2C) across Asia and North America, evaluating UX differentiation to inform AI-product partnership strategy
● Researched U.S. fintech personalization trends and produced investment recommendations calibrated to the Korean market, synthesized into executive reports for the investment committee

PROJECTS

Adaptive Test-Time Sampling for RL-Fine-Tuned Reasoning Policies **Stanford, CA**
Researcher, CS224R: Deep Reinforcement Learning *April 2026 - Present*
● Investigating whether reasoning policies trained with on-policy RL (IPO, RLOO) benefit differently from adaptive test-time compute allocation than SFT-only policies on the Countdown math reasoning task (Qwen 2.5 0.5B base)
● Implementing a self-consistency-based stopping rule that halts sampling once the plurality answer reaches an agreement threshold, redirecting compute toward problems where initial generations disagree
● SFT baseline trained and evaluated (88% test token accuracy, ~71% pass@16); IPO and RLOO checkpoints and the adaptive vs. fixed-k comparison at matched compute budgets in progress

Evaluation User-Style Adaptation for Professional Text Generation **Stanford, CA**
Researcher, CS224N: Natural Language Processing with Deep Learning *January 2026 - March 2026*
● Extended HyPerAlign, a prompt-based personalization framework, with a structured 11-field email schema augmented with user-distinctive vocabulary automatically extracted from writing samples, plus a lightweight prompt optimizer that identifies the minimal effective style field subset
● A 3-field schema (greeting style, punctuation, distinctive voice) closed >50% of the gap to the performance ceiling with a single exemplar, outperforming both parameter-efficient fine-tuning and unstructured hypothesis prompting

Stochastic Routing Process for Cross-Border Payments **Stanford, CA**
Co-Researcher, CS 238: Decision Making Under Uncertainty *October 2025 - December 2025*
● Designed a greedy stochastic routing framework over institution-currency networks; Monte Carlo simulations over 10 years of FX data achieved 100% routing success on feasible corridors
● Final amounts systematically exceeded the commercial mid-market benchmark across all tested corridors; model correctly identified structural failure cases where network asymmetries made routing infeasible

SKILLS & INTERESTS

Programming: Python (PyTorch, HuggingFace Transformers, pandas, NumPy), R, SQL
ML/AI: PyTorch, Hugging Face Transformers, on-policy RL training (IPO, RLOO), supervised fine-tuning, LLM evaluation, Monte Carlo methods
Languages: Spanish (Native), Korean (Native), English (Fluent), Portuguese, Chinese